

# MICROPROCESSOR SYSTEMS

BLG212E

Burak Berk Üstündağ CRN: 11450


Gökhan İnce / Ayşe Yilmazer CRN: 11446

İTÜ Bilgisayar ve Bilişim Fakültesi

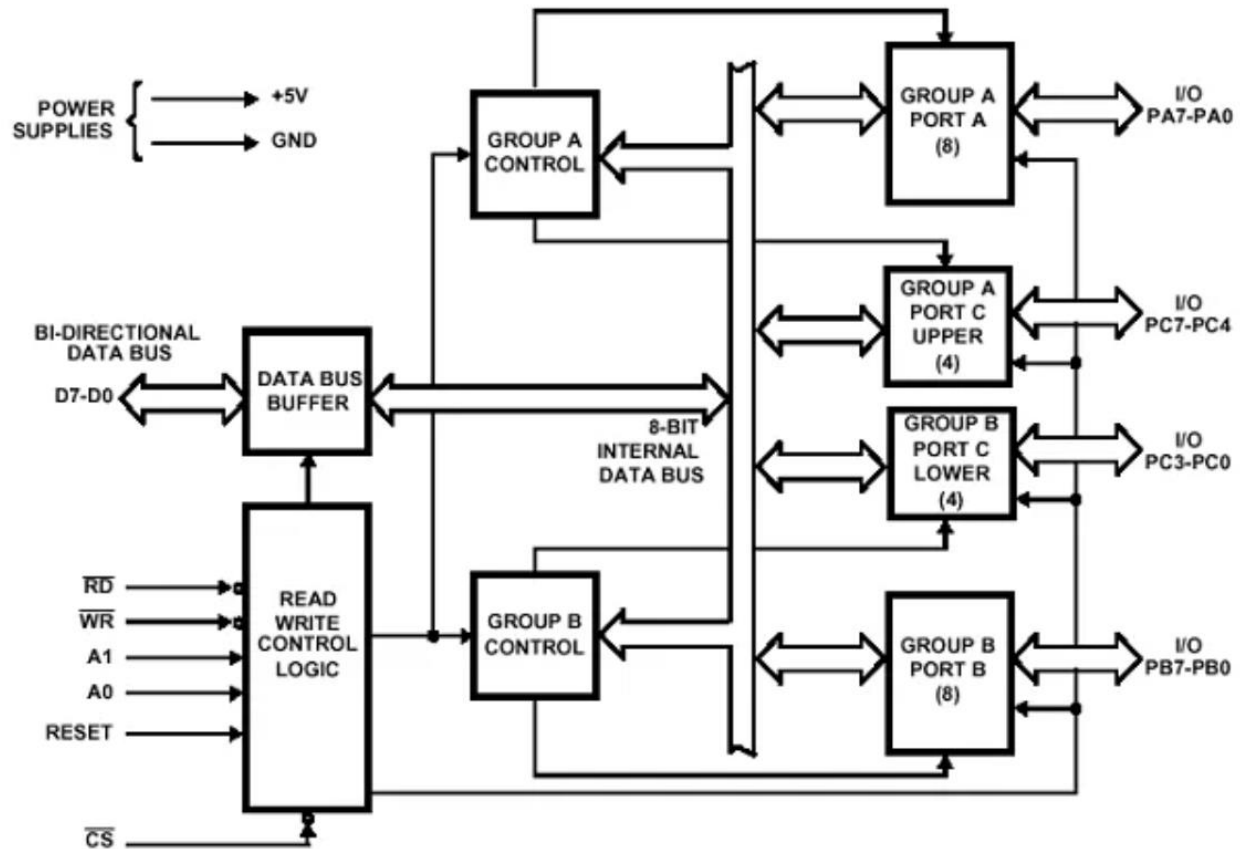
Week 15: Configurable I/O and an introduction to advanced microprocessors



# Configurable I/O Application

## 82C55 CMOS Programmable Peripheral Interface



40 pin Plastic DIP





Russian Clone of 8255

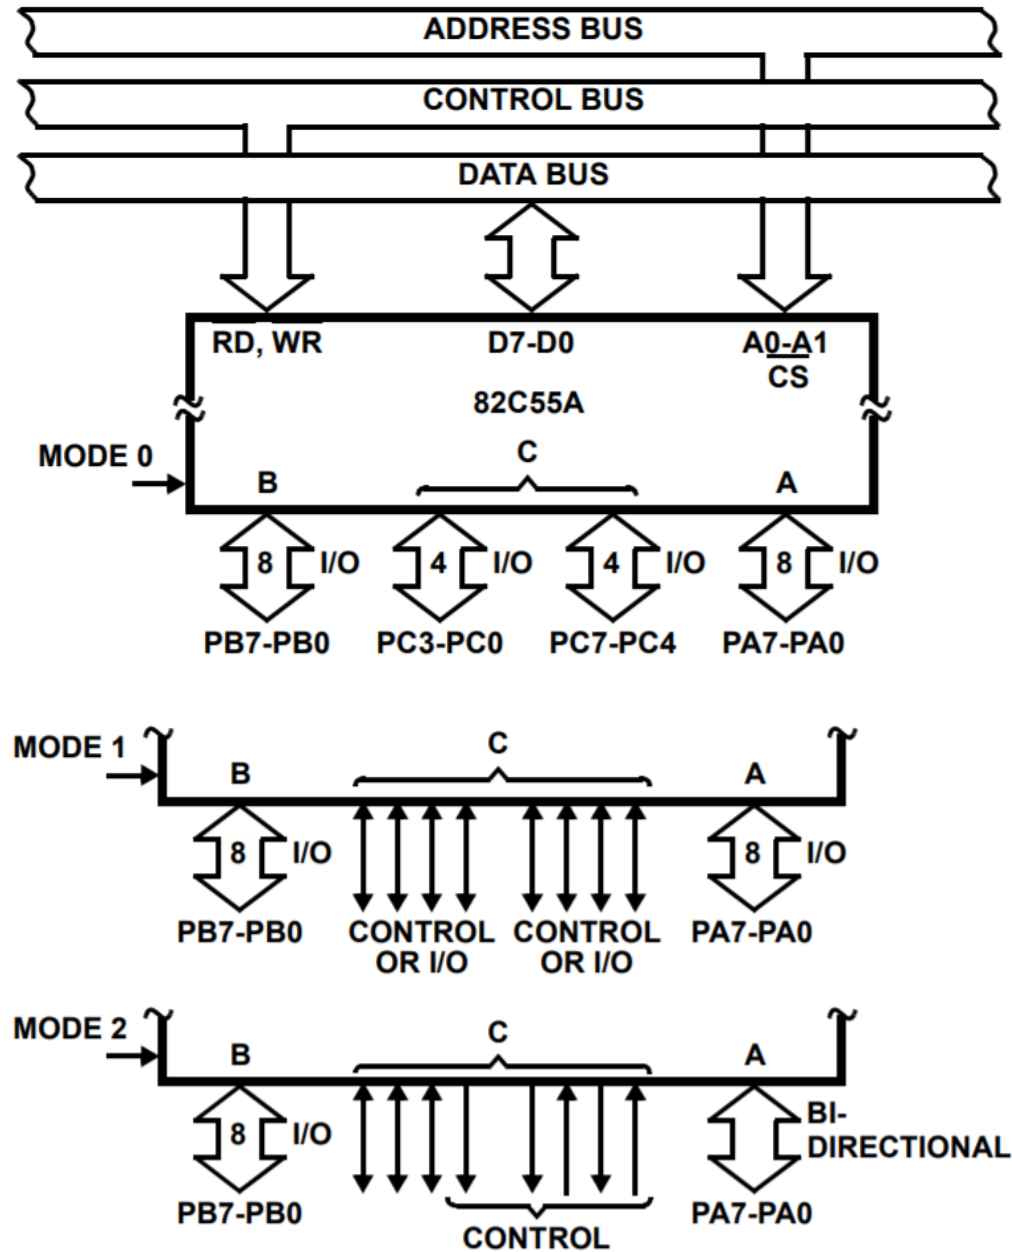
SYMBOL	TYPE	DESCRIPTION
V <sub>CC</sub>		V <sub>CC</sub> : The +5V power supply pin. A 0.1μF capacitor between V <sub>CC</sub> and GND is recommended for decoupling.
GND		GROUND
D0-D7	I/O	DATA BUS: The Data Bus lines are bidirectional three-state pins connected to the system data bus.
RESET	I	RESET: A high on this input clears the control register and all ports (A, B, C) are set to the input mode with the “Bus Hold” circuitry turned on.
$\overline{CS}$	I	CHIP SELECT: Chip select is an active low input used to enable the 82C55A onto the Data Bus for CPU communications.
$\overline{RD}$	I	READ: Read is an active low input control signal used by the CPU to read status information or data via the data bus.
$\overline{WR}$	I	WRITE: Write is an active low input control signal used by the CPU to load control words and data into the 82C55A.
A0-A1	I	ADDRESS: These input signals, in conjunction with the $\overline{RD}$ and $\overline{WR}$ inputs, control the selection of one of the three ports or the control word register. A0 and A1 are normally connected to the least significant bits of the Address Bus A0, A1.
PA0-PA7	I/O	PORT A: 8-bit input and output port. Both bus hold high and bus hold low circuitry are present on this port.
PB0-PB7	I/O	PORT B: 8-bit input and output port. Bus hold high circuitry is present on this port.
PC0-PC7	I/O	PORT C: 8-bit input and output port. Bus hold circuitry is present on this port.

**Basic Operation:**

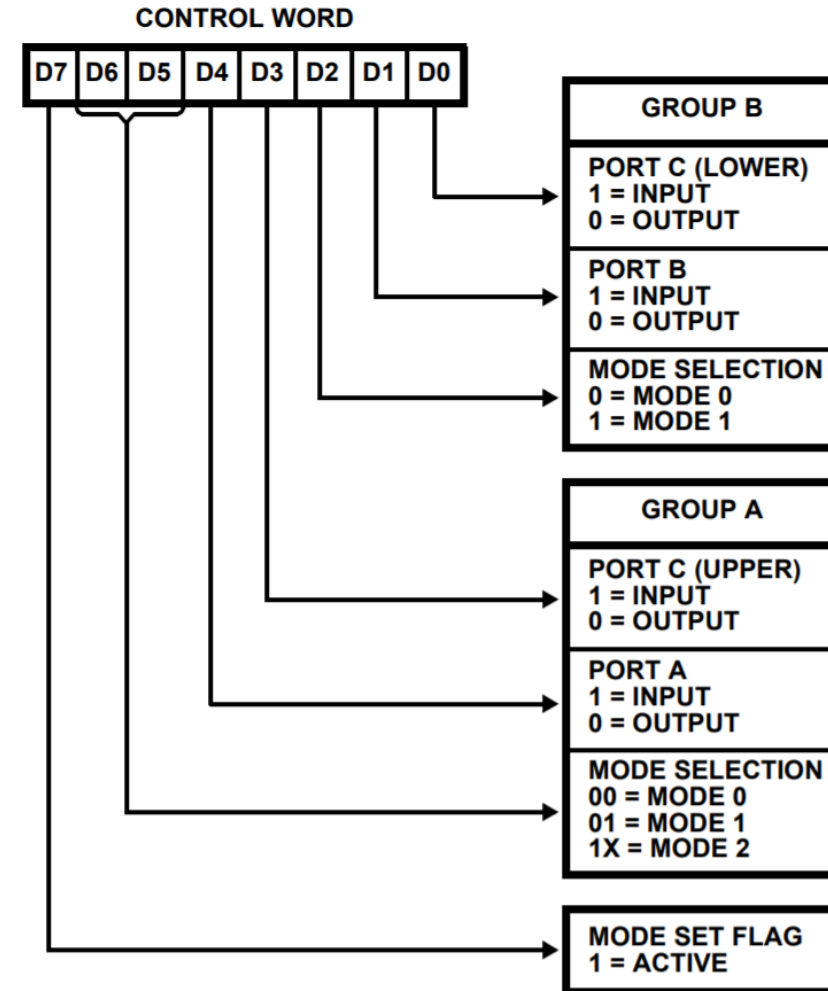
Input Operation (READ)					
A <sub>1</sub>	A <sub>0</sub>	$\overline{RD}$	$\overline{WR}$	$\overline{CS}$	
0	0	0	1	0	PORT A → Data Bus
0	1	0	1	0	PORT B → Data Bus
1	0	0	1	0	PORT C → Data Bus

Output Operation (WRITE)					
A <sub>1</sub>	A <sub>0</sub>	$\overline{RD}$	$\overline{WR}$	$\overline{CS}$	
0	0	1	0	0	Data Bus → PORT A
0	1	1	0	0	Data Bus → PORT B
1	0	1	0	0	Data Bus → PORT C
1	1	1	0	0	Data Bus → CONTROL


Disable Function					
A <sub>1</sub>	A <sub>0</sub>	$\overline{RD}$	$\overline{WR}$	$\overline{CS}$	
X	X	X	X	1	Data Bus → High Z State
X	X	1	1	0	Data Bus → High Z State




## 82C55 Mode Control:



**Mode 0 (Basic Input)**



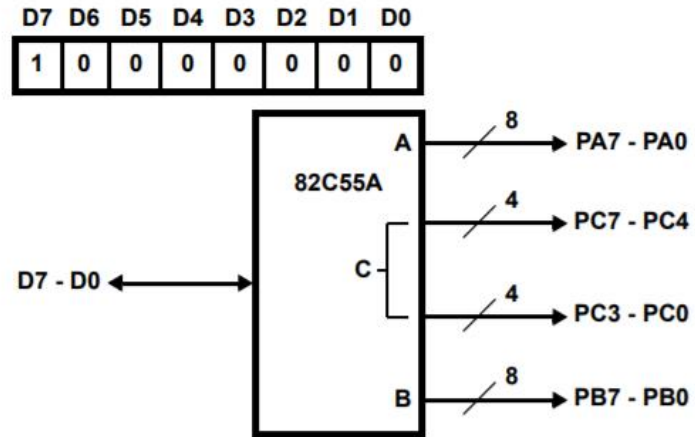
**Mode 0 (Basic Output)**



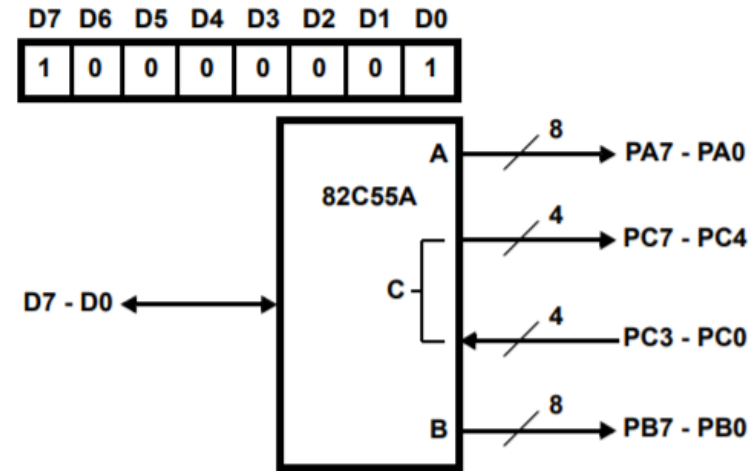
## Control World Examples for 82C55

### Mode 0 Configurations

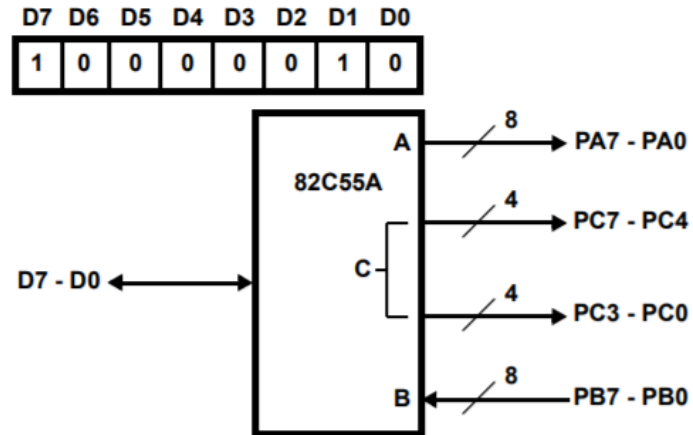
#### CONTROL WORD #0



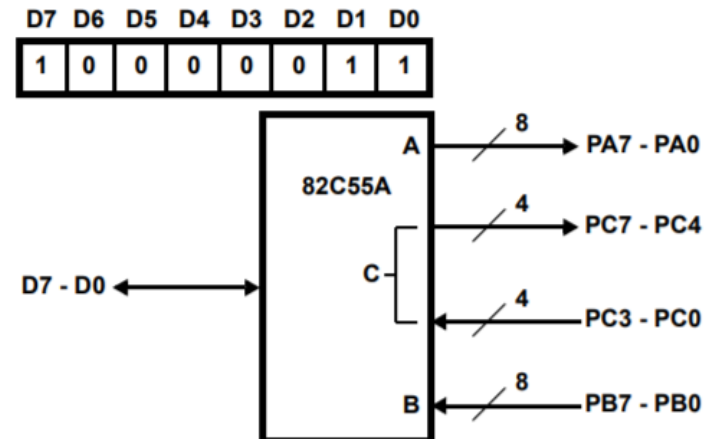
#### CONTROL WORD #1

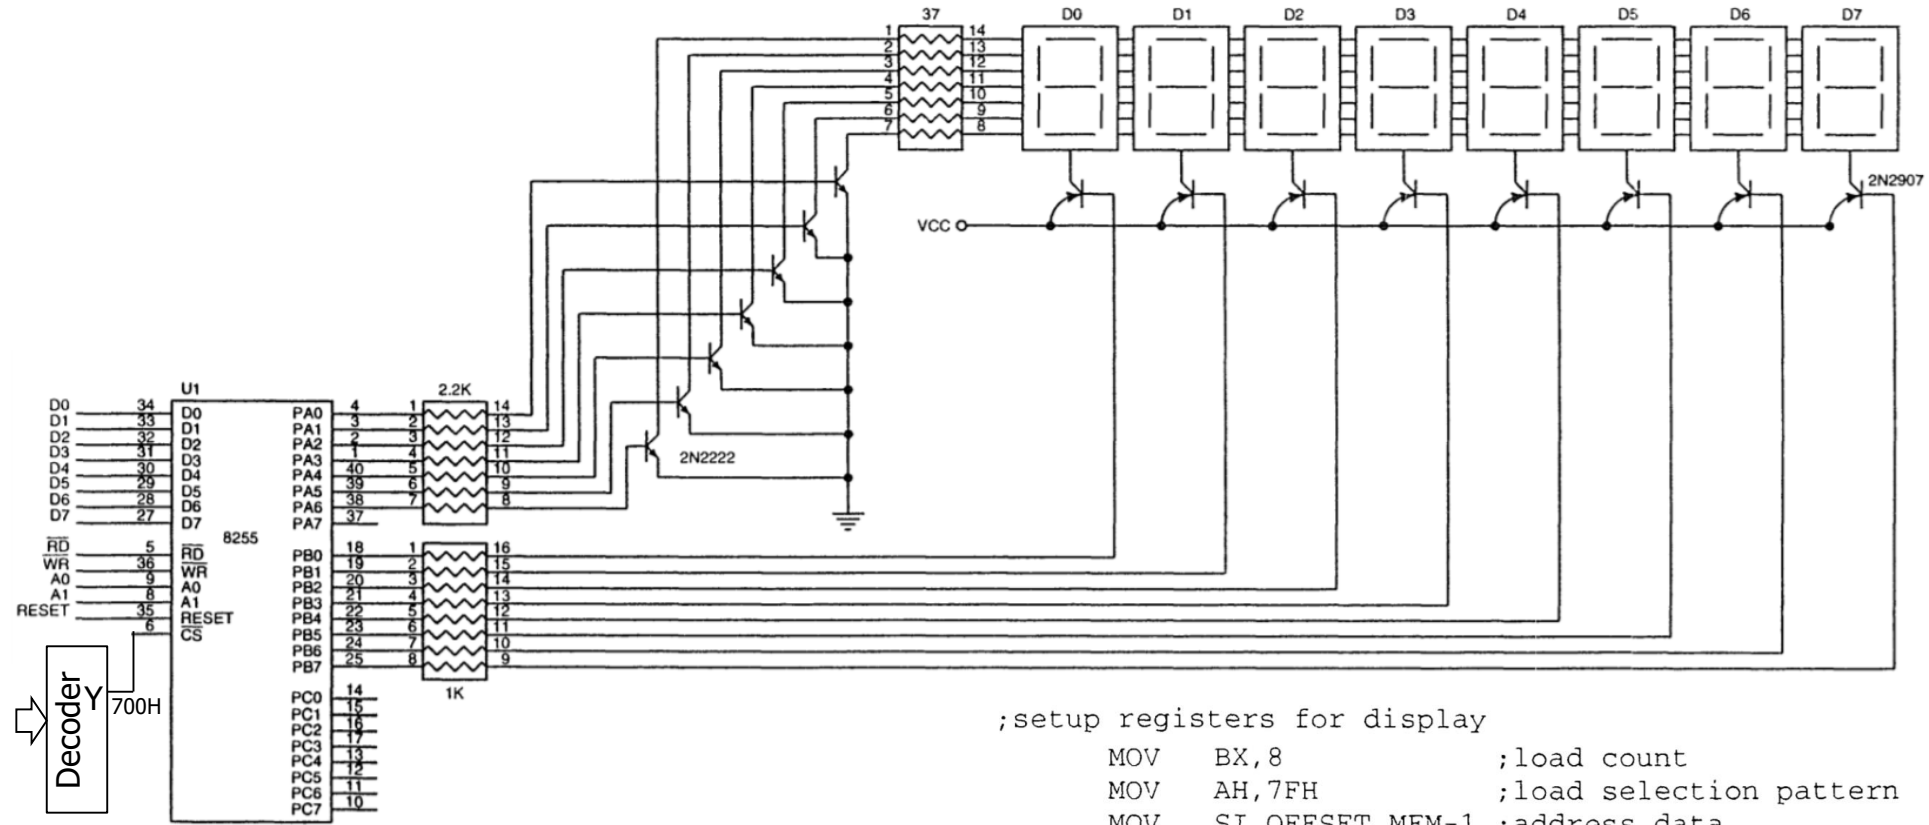


#### CONTROL WORD #2



#### CONTROL WORD #3






Mode 0, output configuration (INTEL):

```
MOV AL,10000000B
MOV DX,703H          ;address command
OUT DX,AL             ;program 82C55
```



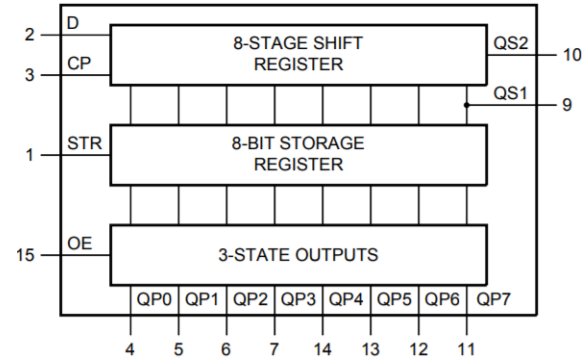
```
;setup registers for display
MOV BX,8              ;load count
MOV AH,7FH            ;load selection pattern
MOV SI,OFFSET MEM-1   ;address data
MOV DX,701H           ;address Port B

;display 8 digits
DISP1:
MOV AL,AH             ;select a digit
OUT DX,AL
DEC DX                ;address Port A
MOV AL,[BX+SI]        ;get 7-segment data
OUT DX,AL
CALL DELAY            ;wait one millisecond
ROR AH,1              ;address next digit
INC DX                ;address Port B
DEC BX                ;adjust count
JNZ DISP1             ;repeat 8 times
```





## Serial Drive with 4094 Shift Register




*H = HIGH voltage level; L = LOW voltage level; X = don't care; Z = HIGH-impedance OFF-state; NC = no change;*


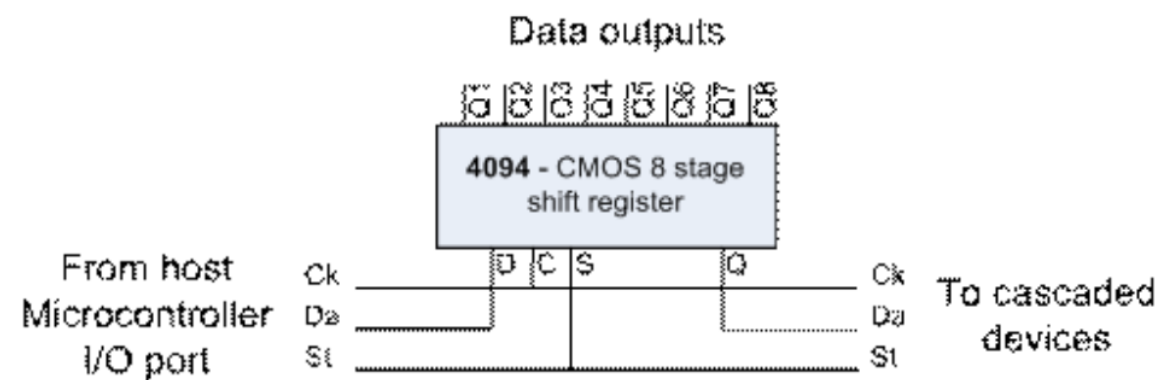
*↑ = positive-going transition; ↓ = negative-going transition;*

*Q6S = the data in register stage 6 before the LOW to HIGH clock transition;*


*Q7S = the data in register stage 7 before the HIGH to LOW clock transition.*

Inputs				Parallel outputs		Serial outputs	
CP	OE	STR	D	QP0	QPn	QS1	QS2
↑	L	X	X	Z	Z	Q6S	NC
↓	L	X	X	Z	Z	NC	Q7S
↑	H	L	X	NC	NC	Q6S	NC
↑	H	H	L	L	QPn -1	Q6S	NC
↑	H	H	H	H	QPn -1	Q6S	NC
↓	H	H	H	NC	NC	NC	Q7S





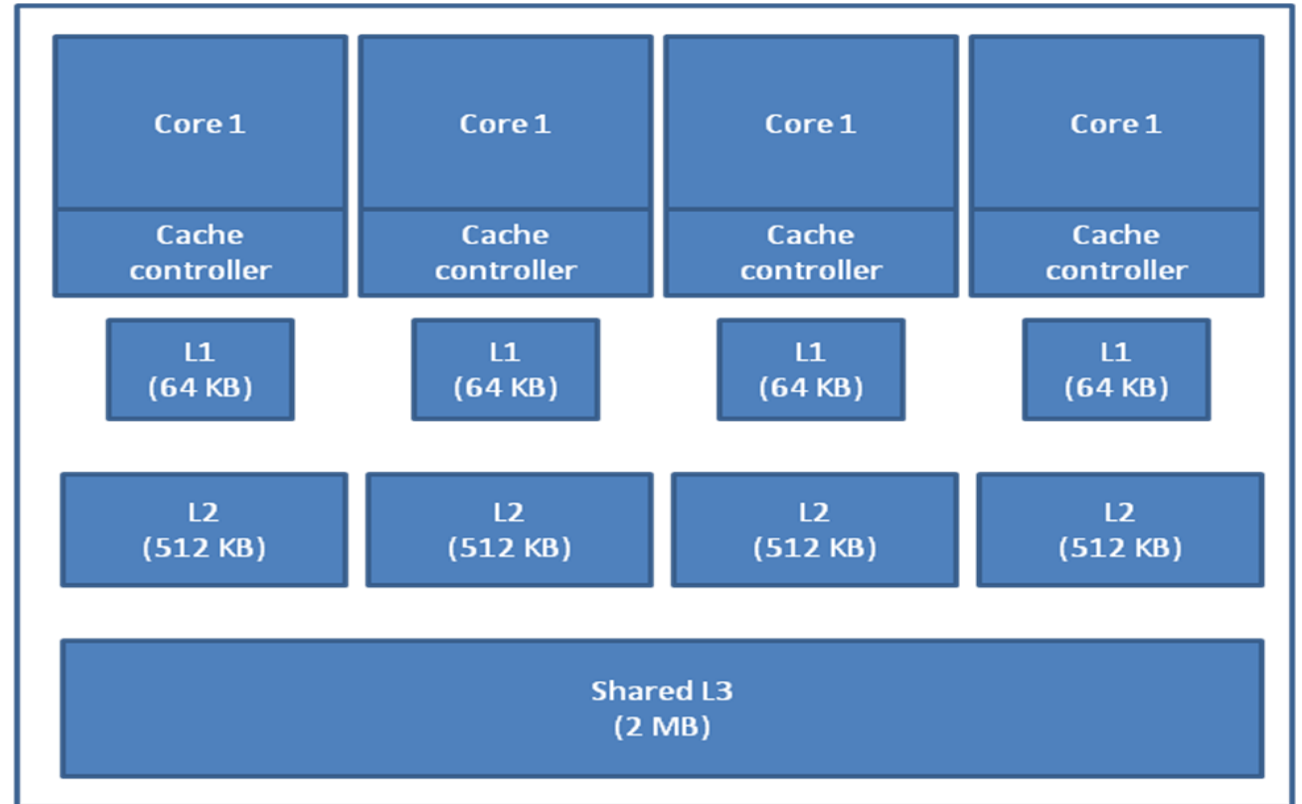
# Memory Hierarchy in Modern Computers



- Processors have cycle times of  $\sim 1$  ns
- Fast DRAM has a cycle time of  $\sim 100$  ns


## Tradeoff:

SRAM is fast  
Economically unfeasible for large memories





Haswell-E die shot The repetitive structures in the middle of the chip are 20MB of shared L3 cache.




# The Concept of a Cache

- Feasible to have small amount of fast memory and/or large amount of slow memory.
- Target:
  - Size advantage of DRAM
  - Speed advantage of SRAM.
- CPU looks in cache for data it seeks from main memory.
- If data not there it retrieves it from main memory.
- If the cache is able to service "most" CPU requests then effectively we will get speed advantage of cache.
- All addresses in cache are also in memory

Increasing speed  
as we get closer to  
the processor

Increasing size as  
we get farther away  
from the processor



# Principle of Locality

- A program tends to access a relatively small region of memory irrespective of its actual memory footprint in any given interval of time. While the region of activity may change over time, such changes are gradual
- Spatial Locality: Tendency for locations close to a location that has been accessed to also be accessed
- Temporal Locality: Tendency for a location that has been accessed to be accessed again
- Example

```
for (i=0; i<100000; i++)  
    a[i] = b[i];
```



# Basic terminologies

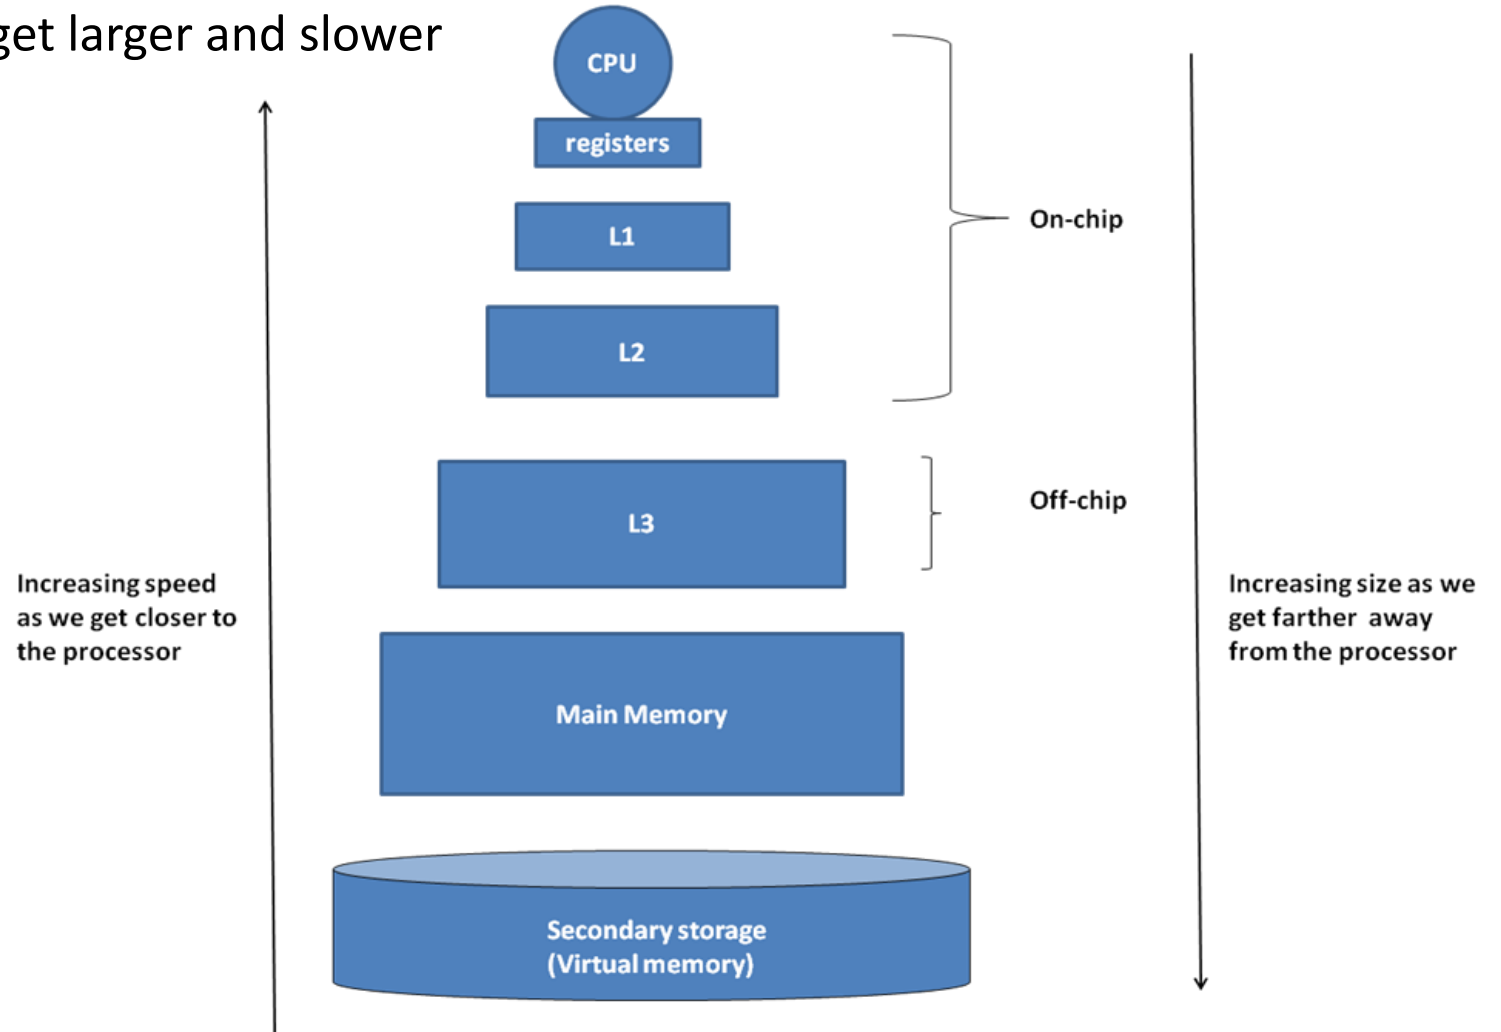
- **Hit**: CPU finding contents of memory address in cache
- **Hit rate** ( $h$ ) is probability of *successful lookup* in cache by CPU.
- **Miss**: CPU *failing* to find what it wants in cache (incurs trip to deeper levels of memory hierarchy)
- **Miss rate** ( $m$ ) is probability of *missing* in cache and is equal to  $1-h$ .
- **Miss penalty**: Time penalty associated with servicing a miss at any particular level of memory hierarchy
- **Effective Memory Access Time (EMAT)**: Effective access time experienced by the CPU when accessing memory.
  - Time to lookup cache to see if memory location is already there
  - Upon cache miss, time to go to deeper levels of memory hierarchy

$$EMAT = T_c + m * T_m$$

where  $m$  is cache miss rate,  $T_c$  the cache access time and  $T_m$  the miss penalty

# Multilevel Memory Hierarchy

- Modern processors use multiple levels of caches.
- As we move away from processor, caches get larger and slower
- $EMAT_i = T_i + m_i * EMAT_{i+1}$
- where  $T_i$  is access time for level  $i$
- and  $m_i$  is miss rate for level  $i$




# Cache organization

- There are three facets to the organization of the cache:
  1. Placement: Where do we place in the cache the data read from the memory?
  2. Algorithm for lookup: How do we find something that we have placed in the cache?
  3. Validity: How do we know if the data in the cache is valid?

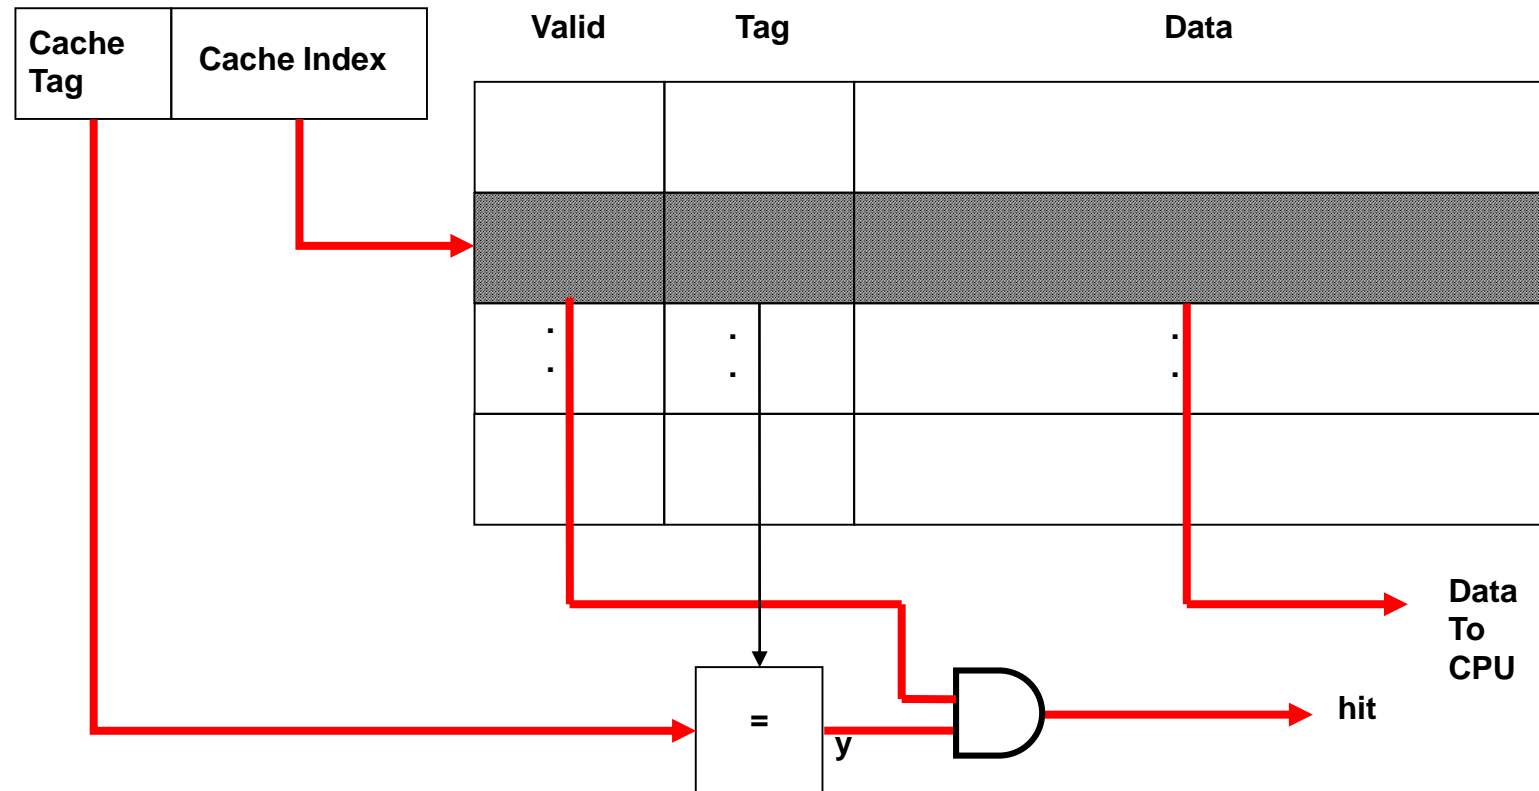
# Cache Lookup

- Example:
  - 4Gb Memory: 32 bit address
  - 256 Kb Cache
  - Cache is organized by words
    - 1 Gword memory
    - 64 Kword cache → 16 bit cache index

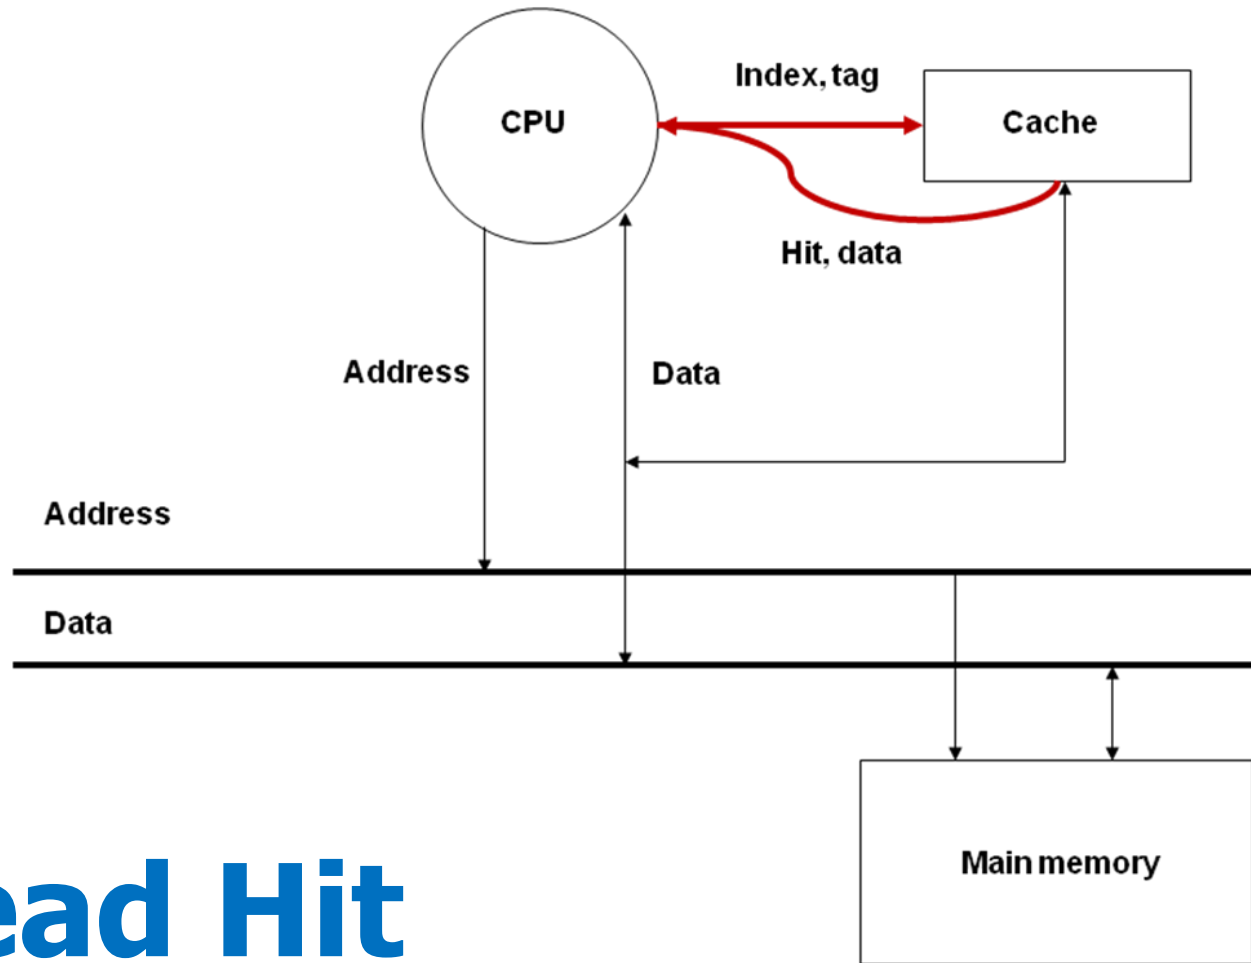


# Hardware for direct mapped cache

Memory address



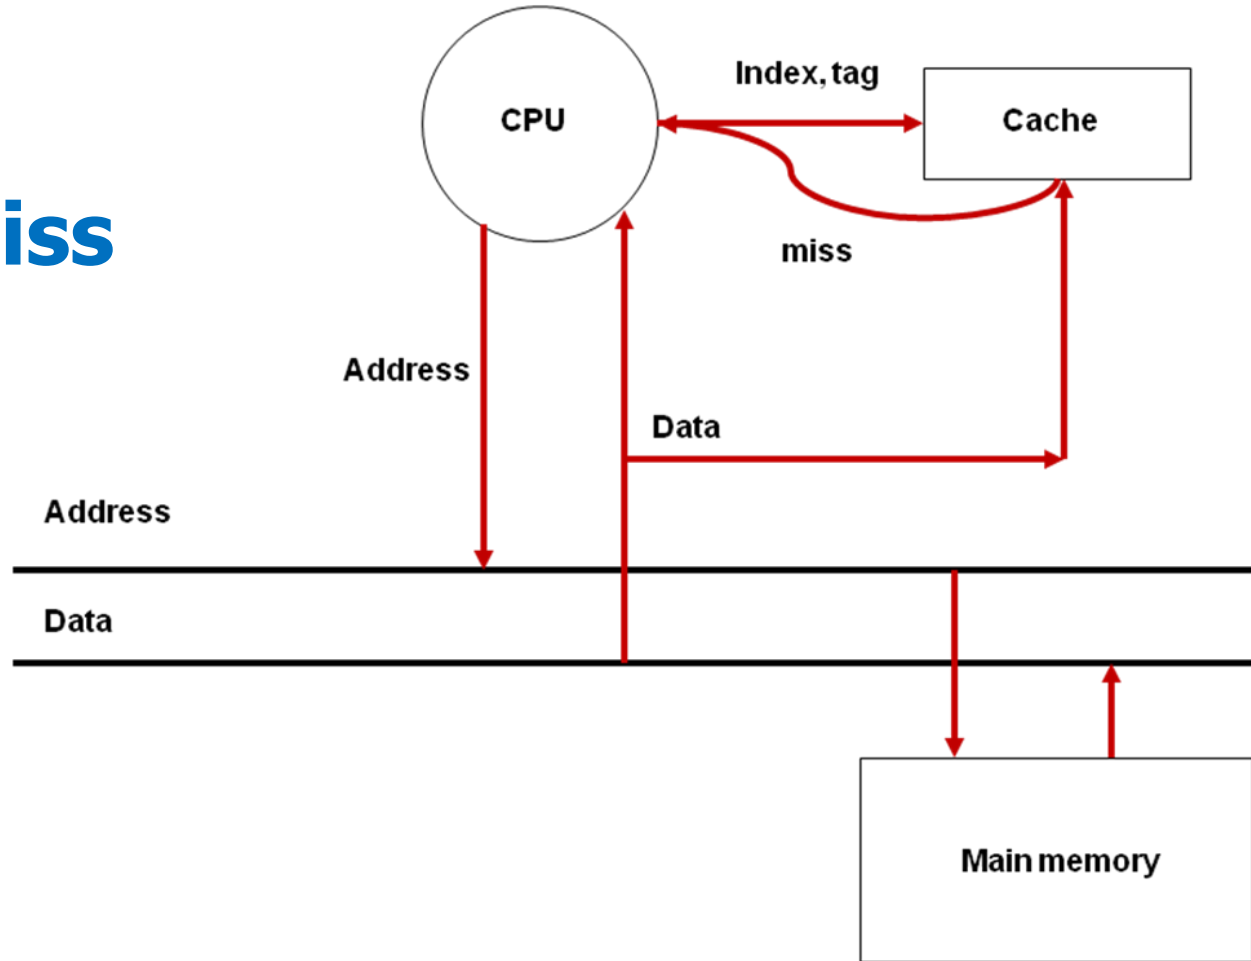
# Cache read/write algorithms



**Read Hit**

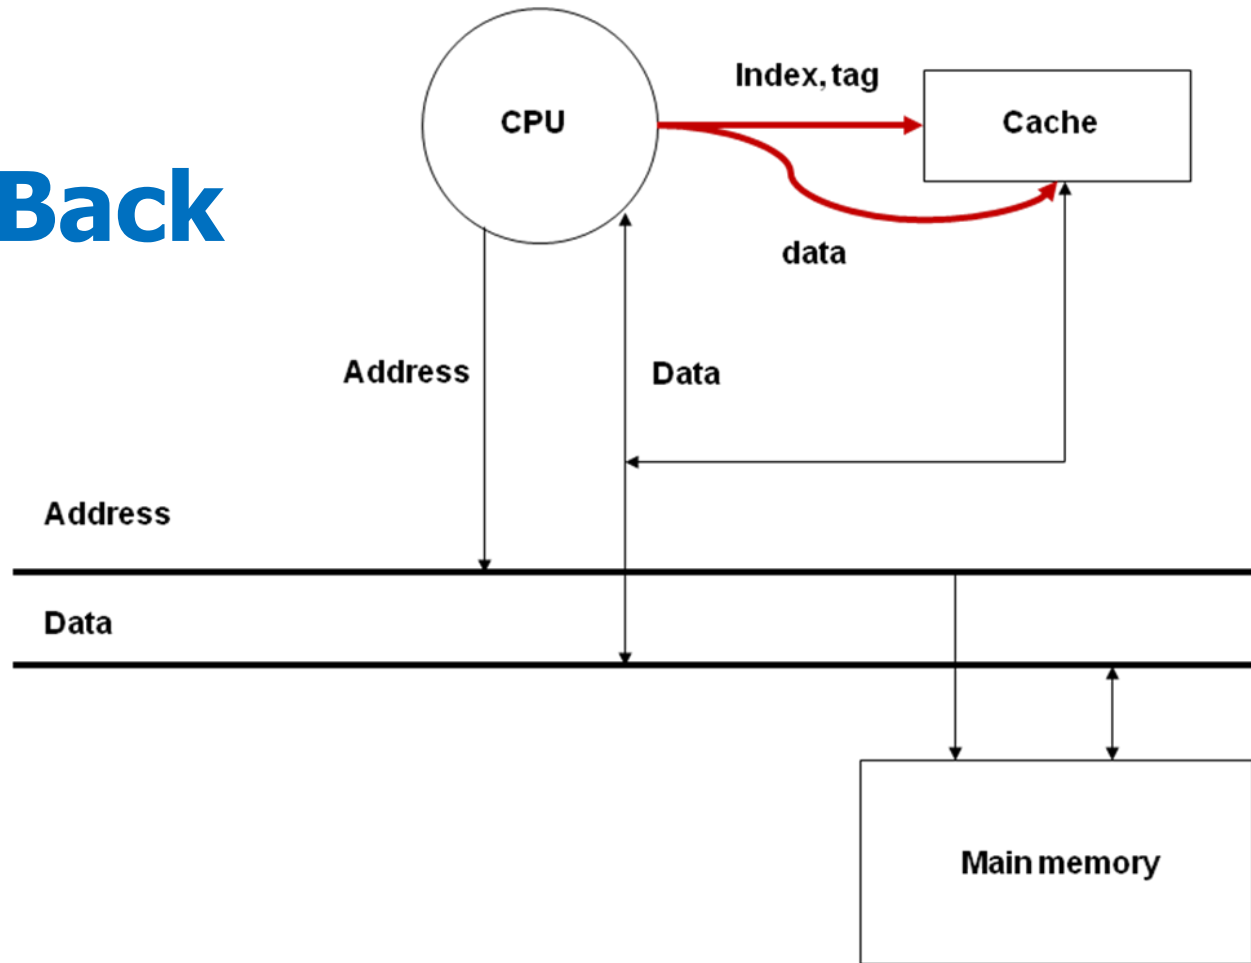
# Basic cache read/write algorithms

## Read Miss



# Basic cache read/write algorithms

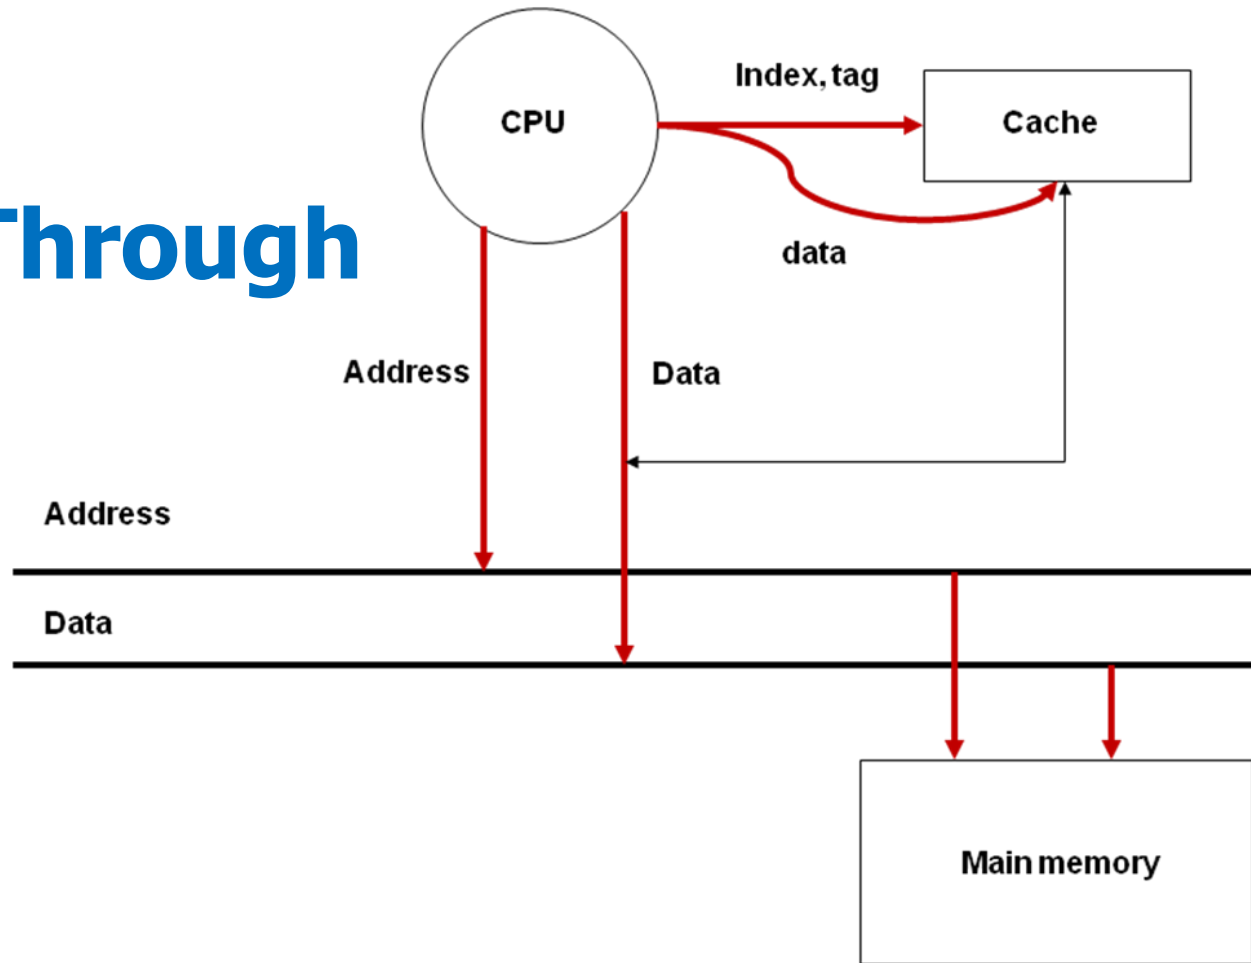
## Write-Back





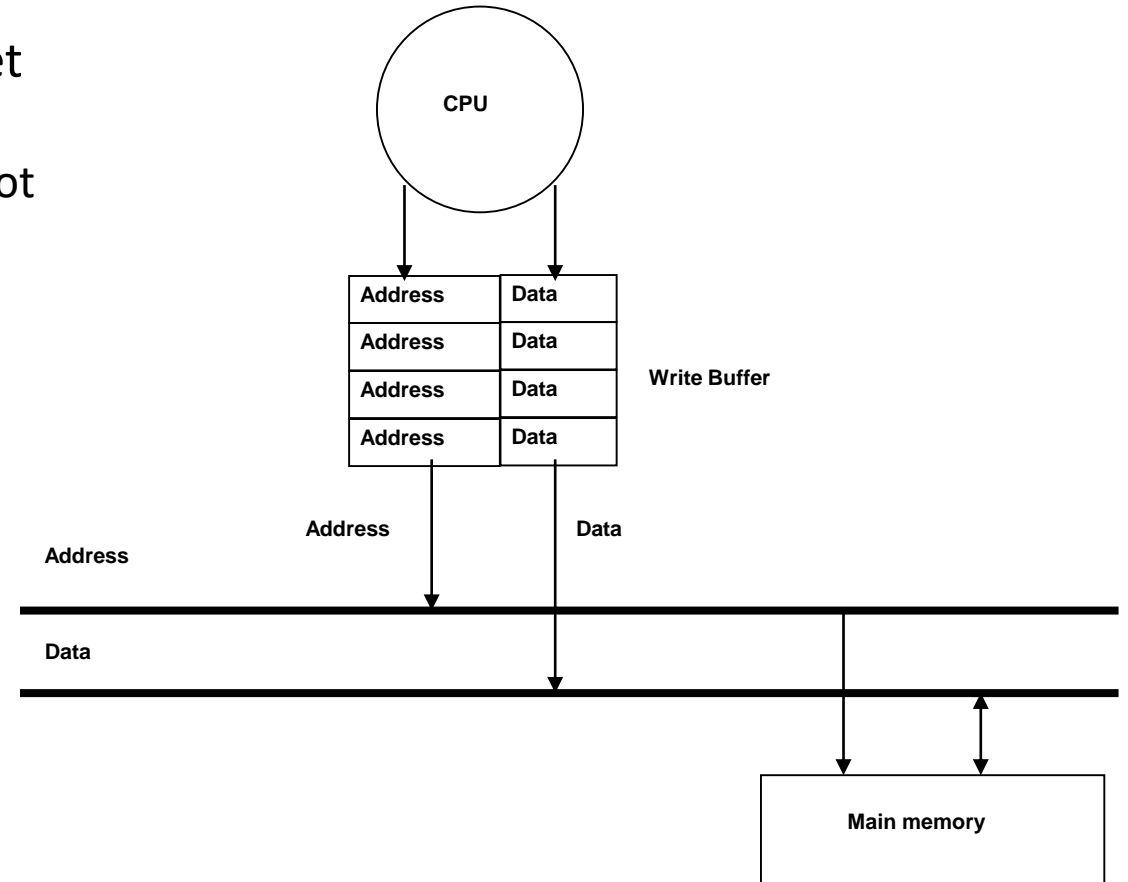
# Basic cache read/write algorithms

## Write-Through



# Write Through Policy

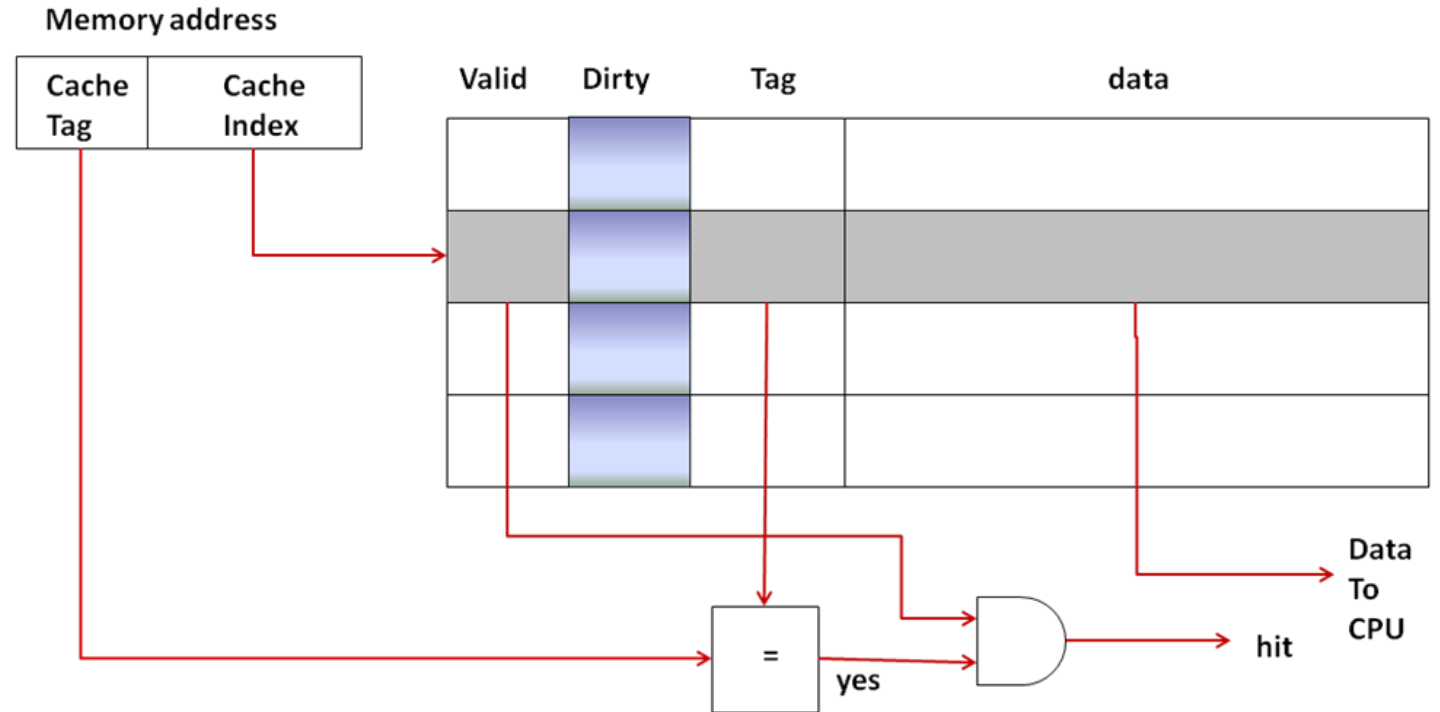
- Each write goes to cache. Tag is set and valid bit is set
  - This is write allocate
  - There is also a no-write allocate where the cache is not written to if there was a write miss
- Each write also goes to write buffer
- Write buffer writes data into main memory
  - Will stall if write buffer full



# Write back policy

- CPU writes data to cache setting dirty bit
  - Note: Cache and memory are now inconsistent but the dirty bit tells us that

- Write Through
  - Cache logic simpler and faster
  - Creates more bus traffic
- Write back
  - Requires dirty bit and extra logic
- Multilevel cache processors may use both
  - L1 Write through
  - L2/L3 Write back



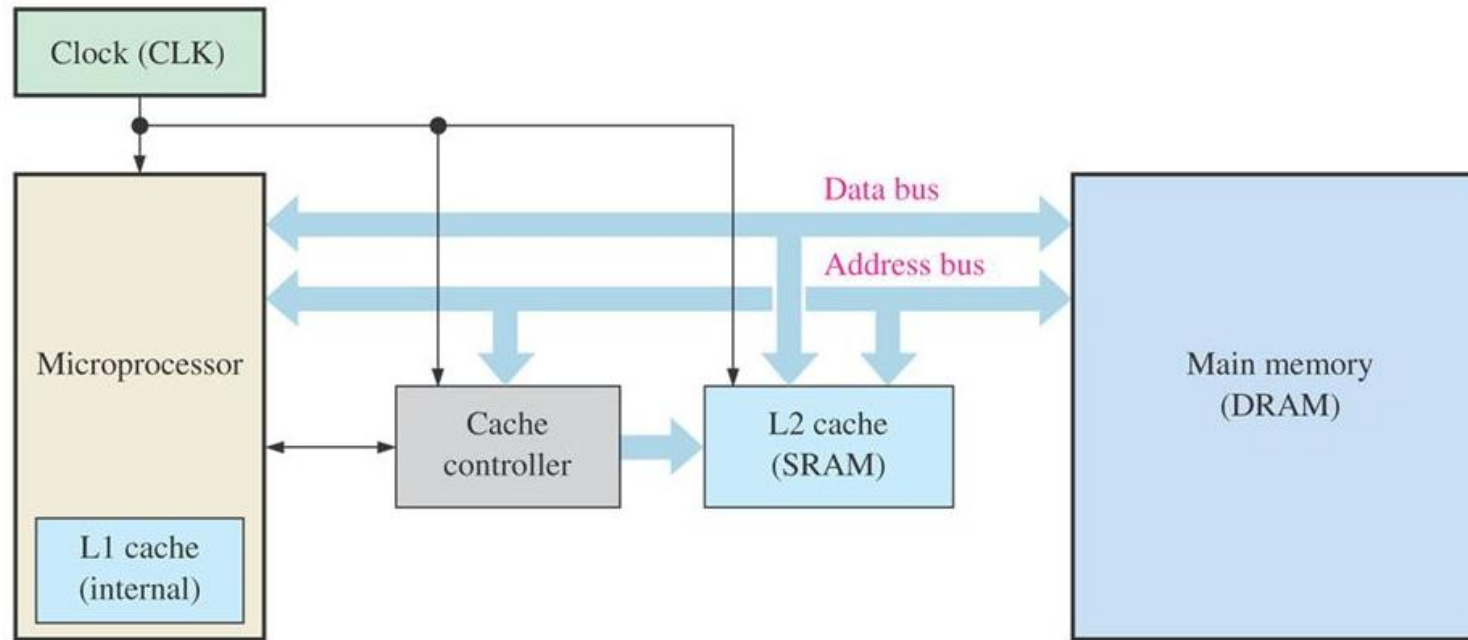
# Cache misses in the processor pipeline

- Read miss in the MEM stage:

```
I1: ld r1, a ; r1 <- MEM[a]
I2: add r3, r4, r5 ; r3 <- r4 + r5
I3: and r6, r7, r8 ; r6 <- r7 AND r8
I4: add r2, r4, r5 ; r2 <- r4 + r5
I5: add r2, r1, r2 ; r2 <- r1 + r2
```

- **Write miss in the MEM stage:** The write-buffer alleviates the ill effects of write misses in the MEM stage. (Write-Through)

# Cache controller



- Upon request from processor, looks up cache to determine hit or miss, serving data up to processor in case of hit.
- Upon miss, initiates bus transaction to read missing block from deeper levels of memory hierarchy.
- Depending on details of memory bus, requested data block may arrive asynchronously with respect to request. In this case, cache controller receives block and places it in appropriate spot in cache.
- Provides ability for the processor to specify certain regions of memory as “uncachable.”

# Cache Design Considerations

- Principles of spatial and temporal locality
- Hit, miss, hit rate, miss rate, cycle time, hit time, miss penalty
- Multilevel caches and design considerations thereof
- Direct mapped caches
- Cache read/write algorithms
- Spatial locality and blocksize
- Fully- and set-associative caches
- Considerations for I- and D-caches
- Cache replacement policy
- Types of misses
- TLB and caches
- Cache controller
- Virtually indexed physically tagged caches

Category	Vocabulary	Details
Principle of locality	Spatial	Access to contiguous memory locations
	Temporal	Reuse of memory locations already accessed
Cache organization	Direct-mapped	One-to-one mapping
	Fully associative	One-to-any mapping
	Set associative	One-to-many mapping
Cache reading/writing	Read hit/Write hit	Memory location being accessed by the CPU is present in the cache
	Read miss/Write miss	Memory location being accessed by the CPU is not present in the cache
Cache write policy	Write through	CPU writes to cache and memory
	Write back	CPU only writes to cache; memory updated on replacement
Cache parameters	Total cache size ( $S$ )	Total data size of cache in bytes
	Block Size ( $B$ )	Size of contiguous data in one data block
	Degree of associativity ( $p$ )	Number of homes a given memory block can reside in a cache
	Number of cache lines ( $L$ )	$S/pB$
	Cache access time	Time in CPU clock cycles to check hit/miss in cache
	Unit of CPU access	Size of data exchange between CPU and cache
	Unit of memory transfer	Size of data exchange between cache and memory
	Miss penalty	Time in CPU clock cycles to handle a cache miss
Memory address interpretation	Index ( $n$ )	$\log_2 L$ bits, used to look up a particular cache line
	Block offset ( $b$ )	$\log_2 B$ bits, used to select a specific byte within a block
	Tag ( $t$ )	$a - (n+b)$ bits, where $a$ is number of bits in memory address; used for matching with tag stored in the cache

## An example memory hierarchy of modern microprocessors

Type of Memory	Typical Size	Approximate latency in CPU clock cycles to read one word of 4 bytes
CPU registers	8 to 32	Usually immediate access (0-1 clock cycles)
L1 Cache	32 (Kilobyte) KB to 128 KB	3 clock cycles
L2 Cache	128 KB to 4 Megabyte (MB)	10 clock cycles
Main (Physical) Memory	256 MB to 4 Gigabyte (GB)	100 clock cycles
Virtual Memory (on disk)	1 GB to 1 Terabyte (TB)	1000 to 10,000 clock cycles (not accounting for the software overhead of handling page faults)



Cache	Bulldozer	Piledriver	Steamroller
Level 1 code	64 kB, 2-way, 64 B line size, shared between two cores.	64 kB, 2-way, 64 B line size, shared between two cores.	96 kB, 3-way, 64 B line size, shared between two cores.
Level 1 data	16 kB, 4-way, 64 B line size, per core. Latency 3-4 clocks.	16 kB, 4-way, 64 B line size, per core. Latency 3-4 clocks.	16 kB, 4-way, 64 B line size, per core. Latency 3-4 clocks.
Level 2	1 - 2 MB, 16-way, 64 B line size, shared between two cores. Latency 21 clocks. Read throughput 1 per 4 clock. Write throughput 1 per 12 clock.	2 MB, 16-way, 64 B line size, shared between two cores. Latency 20 clocks. Read throughput 1 per 4 clock. Write throughput 1 per 12 clock.	2 MB, 16-way, 64 B line size, shared between two cores. Latency 19 clocks. Read throughput 1 per 4 clock. Write throughput 1 per 6 clock.
Level 3	0 - 8 MB, 64-way, 64 B line size, shared between all cores. Latency 87 clock. Read throughput 1 per 15 clock. Write throughput 1 per 21 clock.	0 - 8 MB, 64-way, 64 B line size, shared between all cores. Latency 87 clock. Read throughput 1 per 15 clock. Write throughput 1 per 21 clock.	None
<b>Table      Cache sizes on AMD Bulldozer, Piledriver and Steamroller</b>			